

Python for Data Data for Python

Barry Rowlingson



Lancaster University
Medical School

**This is all about reading and writing data in
Python...**

Reading Text Files with open

```
This is the first message.  
The second message is this.  
Third message.
```

The simplest reading we ever need to do is to read some lines from a plain text file.
Here's the file.

Reading Text Files with open

```
>>> h = open("messages.txt","r")  
>>> h  
<_io.TextIOWrapper name='messages.txt' mode='r' encoding='UTF-8'>
```

Open the file for reading

Because reading a file is asking the computer for some resources, you can't dive straight in and get the data. It's good manners to tell the operating system you are interested in this data file, so most modern languages require you to "open" the file for reading before you start with it. If the file doesn't exist or there's some other problem, the OS will let Python know and there will be an error at this point.

Reading Text Files with open

```
>>> h = open("messages.txt","r")
>>> h
<_io.TextIOWrapper name='messages.txt' mode='r' encoding='UTF-8'>
```

h is... one of these.
A connection or a
handle

In Python once you've opened the file you get back one of these objects, referred to as a "handle" or a "connection" or something simply a "file" object.

Reading Text Files with open

```
>>> h = open("messages.txt","r")
>>> h
<_io.TextIOWrapper name='messages.txt' mode='r' encoding='UTF-8'>

>>> h.readline()
'This is the first message.\n'
```

Read the first line...

Now you can read a line with the “readline” method, and it is returned.

Reading Text Files with open

```
>>> h = open("messages.txt","r")  
>>> h  
<_io.TextIOWrapper name='messages.txt' mode='r' encoding='UTF-8'>
```

```
>>> h.readline()  
'This is the first message.\n'
```

Read the first line...

```
>>> h.readline().strip()  
'The second message is this.'
```

Read the second line
but strip line ending

Notice the backslash-N on the string? That's the end-of-line marker. If you don't want that (and you usually don't) you can use the "strip" method of text strings to remove it. Notice that when I do a "readline" the second time, it's now reading the second line.

Reading Text Files with open

```
>>> h = open("messages.txt", "r")
>>> h
<_io.TextIOWrapper name='messages.txt' mode='r' encoding='UTF-8'>

>>> h.readline()
'This is the first message.\n'
>>> h.readline().strip()
'The second message is this.'


>>> h.close()
>>> h.readline()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: I/O operation on closed file.
```

Close the connection.
No more reading.

Once you've finished reading, it's good practice to "close" the connection. Most operating systems have a limit on the number of open files per process. Python will close them automatically when an open goes out of scope, but if you should always remember this is happening. Operations on a closed file handle will fail.

Read a whole file

```
>>> h = open("messages.txt")
>>> h.readlines()
['This is the first message.\n', 'The second message is this.\n',
 'Third message.\n']
```



Get a list of all lines with
readlines

If you want to slurp in every line of a file, then there's the "readlines" method. You get back a list of all the lines, complete with the line ending characters.

Read a whole file

```
>>> h = open("messages.txt")
>>> h.readlines()
['This is the first message.\n', 'The second message is this.\n',
 'Third message.\n']
```

```
>>> h.seek(0)
0
```

Rewind...

Iterate over lines, stripping
end of line character

```
>>> [line.strip() for line in h.readlines()]
['This is the first message.', 'The second message is this.',
 'Third message.']
```

Let's strip those off. First I use the "seek" method to rewind the file back to the start, and then I use a list comprehension that calls "strip" on every line it reads. Now I get a list of lines without end-of-line characters.

Read a whole file

```
>>> h = open("messages.txt")
>>> h.readlines()
['This is the first message.\n', 'The second message is this.\n',
 'Third message.\n']
```

```
>>> h.seek(0)
0
>>> [line.strip() for line in h.readlines()]
['This is the first message.', 'The second message is this.',
 'Third message.']
```

Can iterate over the handle

```
>>> h = open("messages.txt")
>>> [line.strip() for line in h]
['This is the first message.', 'The second message is this.',
 'Third message.']
```

I can do that even quicker by iterating over the handle - Python understands that an iterative operation over a file handle is likely to be iterating over reading lines, so there's a little shortcut that lets you do this.

Or even shorter...

The list function iterates:

```
>>> list(h)
['This is the first message.\n', 'The second message is this.\n', 'Third messa
```

One-liner:


```
>>> L = list(open("messages.txt"))
>>> L
['This is the first message.\n', 'The second message is this.\n', 'Third messa
```

...brevity is not always best practice!

You can make this shorter (without the strip method) by noticing that the “list” function iterates over its argument, and so a simple call to “list” on the file handle will return all the lines. You can then read all the lines with a one-liner, although its not always best practice to cut the number of lines to a minimum.

Why would you read lines?

```
3
4
0,0
1,0
1,1
0,1
3
1,1
0.5,2
0,1
5
0.5, 0.1
0.7, 0.3
0.6, 0.4
0.4, 0.3
0.3, 0.1
```



Why would you want to read a line at a time? Here's a file format that describes a set of shapes.

File format

```
3
4
0,0
1,0
1,1
0,1
3
1,1
0.5,2
0,1
5
0.5, 0.1
0.7, 0.3
0.6, 0.4
0.4, 0.3
0.3, 0.1
```


Contains 3 shapes

First shape has 4 points

Here's 4 x,y, points

Next shape has 3 points

Last shape has 5 points



The first line is the number of shapes. Then for each shape there's the number of points in the shape, followed by that number of lines with the xy coordinates of the shape. Then the next shape starts.

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Here's a function that reads in and prints out that shape data.

Reading

Open the file for reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

First we open the file for reading...

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Read the first line,
convert to integer

Then the next line is the number of shapes, we convert that to an integer.

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Loop over the number of
shapes

Next we loop over the number of shapes....

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Read the number of
points for this shape

The next line is the number of points for this shape. Again we convert to integer.

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Loop for that number of points...

Now we have an inner loop for the number of points...

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Get the line, strip any
space, split on comma

Read the line, strip the line-end marker, split the text by comma. This gets back two text values into x and y

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Convert to numeric format

These need to be decimal numbers so we convert them with "float".

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Print out the values

And then we can print them out, here using the very handy f-string notation.

Reading

```
def readshapes(path):  
    shapedata = open(path, "r")  
    number_of_shapes = int(shapedata.readline())  
    for shape in range(number_of_shapes):  
        number_of_points = int(shapedata.readline())  
        for point in range(number_of_points):  
            x,y = shapedata.readline().strip().split(",")  
            x,y = float(x), float(y)  
            print(f"{shape=} {x=}, {y=}")  
    shapedata.close()
```

Clean up

And finally close the connection.

Run!

```
shape=0 x=0.0, y=0.0  
shape=0 x=1.0, y=0.0  
shape=0 x=1.0, y=1.0  
shape=0 x=0.0, y=1.0  
shape=1 x=1.0, y=1.0  
shape=1 x=0.5, y=2.0  
shape=1 x=0.0, y=1.0  
shape=2 x=0.5, y=0.1  
shape=2 x=0.7, y=0.3  
shape=2 x=0.6, y=0.4  
shape=2 x=0.4, y=0.3  
shape=2 x=0.3, y=0.1
```

Run that, and this gets printed. How would you return these values instead of printing them out?

Writing Text Files

```
>>> newdata = open("newfile.txt", "w")
>>> newdata.write("Hello\n")
6
>>> newdata.write("This is a test\n")
15
>>> newdata.close()
```

Open file for writing

Write a couple of lines

Closing is **very** important

```
>>> newdata = open("newfile.txt", "a")
>>> newdata.write("This is line 3\n")
15
>>> newdata.close()
```

Open file for **appending**

File now has 3 lines

Writing text files is similar to reading them except you use the “write” method instead of “read” ones. Note that closing the file is very important here since the operating system might not have actually written the data to disk until the “close” operation is complete. You can also open a file for appending lines if you want to add text to a file.

Reading CSV Files

```
"x","y","cc"  
359014,416976,0  
352909,426935,0  
353848,422172,0  
359202,417326,0  
357795,415825,0  
352784,426890,0
```

```
def read_casecontrol_list(path):  
    csvfile = open(path)  
    head = csvfile.readline()  
    rows = []  
    for row in csvfile:  
        rows.append(row.strip().split(","))  
    return rows
```

Open, read rows, split on
comma, build a list

```
[['359014', '416976', '0'], ['352909', '426935', '0'], ['353848', '422172', '0'], ['359202', '417326', '0'], ['357795', '415825', '0'], ['352784', '426890', '0']]
```

List of rows, as character

Commonly our data is a bit more structured, often as a tabular structure in a CSV file. You could open the file and read the lines in using a loop, splitting on commas to get text values for each row.

Reading CSV Files

Use the `csv` module

```
import csv

def read_casecontrol(path):
    with open(path) as csvfile:
        ccreader = csv.reader(csvfile)
        header = next(ccreader)
        return [row for row in ccreader]
```

```
[['359014', '416976', '0'], ['352909', '426935', '0'], ['353848', '422172', '0'], ['35920
```

But this is such a common thing that there is a built-in python module for reading CSV files. Import it...

Reading CSV Files

Use the `csv` module

```
import csv
```

Open the file in a `with`
block

```
def read_casecontrol(path):  
    with open(path) as csvfile:  
        ccreader = csv.reader(csvfile)  
        header = next(ccreader)  
        return [row for row in ccreader]
```

```
[['359014', '416976', '0'], ['352909', '426935', '0'], ['353848', '422172', '0'], ['35920
```

Then open the file in a “with” block. This will ensure the file is closed at the end of the code.

Reading CSV Files

```
import csv

def read_casecontrol(path):
    with open(path) as csvfile:
        ccreader = csv.reader(csvfile)
        header = next(ccreader)
        return [row for row in ccreader]
```

Use the `csv` module

Open the file in a `with` block

Get a `csv.reader` object

```
[['359014', '416976', '0'], ['352909', '426935', '0'], ['353848', '422172', '0'], ['35920
```

Then we get something like the ordinary file handle, but it knows its reading a CSV...

Reading CSV Files

```
import csv

def read_casecontrol(path):
    with open(path) as csvfile:
        ccreader = csv.reader(csvfile)
        header = next(ccreader)
        return [row for row in ccreader]
```

Use the `csv` module

Open the file in a `with` block

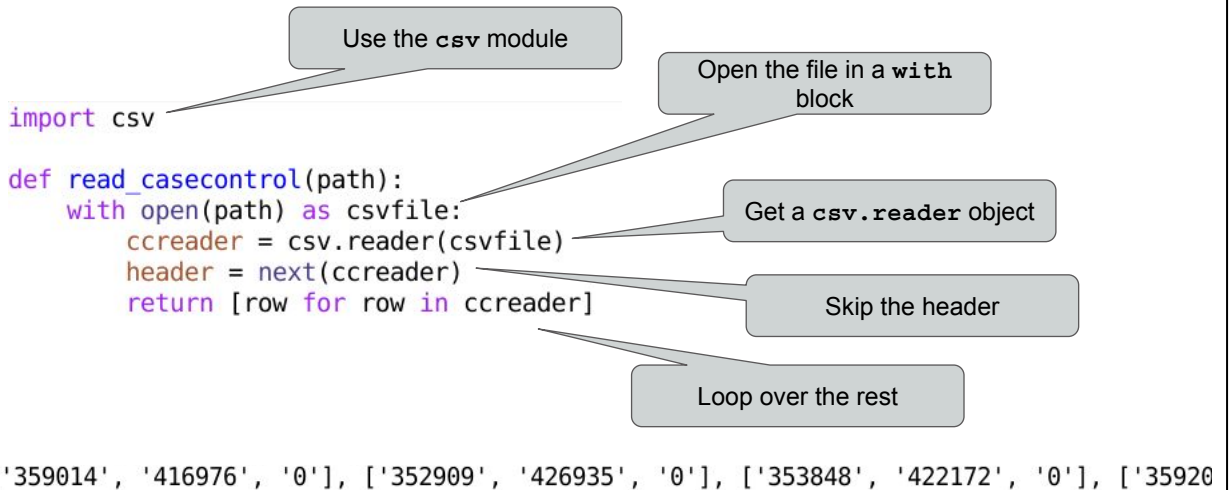
Get a `csv.reader` object

Skip the header

```
[['359014', '416976', '0'], ['352909', '426935', '0'], ['353848', '422172', '0'], ['35920
```

That object is “iterable”, so we can call “next” on it to get the next line. This gets us the headers.

Reading CSV Files



Then we can iterate to get the rest of the lines, and you can see them as a list of lists of character-based items.

Inferring separator, etc

```
import csv

def read_sniff(path):
    with open(path, newline='') as csvfile:
        dialect = csv.Sniffer().sniff(csvfile.read(1024))
        csvfile.seek(0)
        reader = csv.reader(csvfile, dialect)
        return [ row for row in reader ]
```

Have a guess at the "dialect"
from the first 1024 bytes

Read using that dialect.

CSV files can have different separators (commas, semicolons, spaces etc) or other oddities that we can think of as different "dialects" of CSV. The csv package has a mechanism for having a guess at the dialect using a "sniffer". This looks at the first 1024 characters in the file and if it sees commas, it sets the separator to comma, and so on. Then you can seek back to the start of the file and read the rows.

Side Note: Incomprehension

Note that:

```
return [ row for row in reader ]
```

Is nearly always better expressed as:

```
return list(reader)
```

Use that structure if modifying (or filtering) the elements:

```
[line.strip() for line in h]
```

Note I've done a non-pythonic thing in the last couple of slides. This list comprehension is really a sign that I've forgotten that I can just iterate over an iterable with the "list" function. I only really need a comprehension if I want to modify or filter the elements in a list.

Writing CSV Files

```
def write_csv(path, datalist):  
    with open(path, "w") as f:  
        writer = csv.writer(f)  
        writer.writerows(datalist)
```

```
>>> write_csv("out.csv", [ ["Name", "Age"], ["Al", 22], ["Bob", 44], ["Chip", 33] ])
```

```
Name, Age  
Al, 22  
Bob, 44  
Chip, 33
```

Writing data to a CSV is similar. You can take a list of lists and send them to a file with a csv writer object.

Data Frames with Pandas

<https://pandas.pydata.org/>

- Data Series
- Data Frame

Part of “PyData”

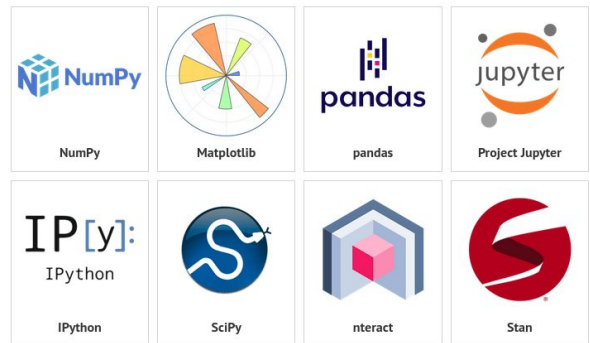


Python's lists and dictionaries are okay for storing simple data, but for more complex situations we need a more capable data structure. And for tabular data we have the “pandas” project, which gives us data series and data frames, as well as some other data types. Its part of the PyData project.

numfocus.org



NUMFOCUS
OPEN CODE = BETTER SCIENCE



....and many more

PyData is a project of the NumFocus organisation, a US charity that sponsors various open source projects in the data science arena. Some we've used already.

-
- `import pandas as pd`
 - A table is a **DataFrame**
 - Each column is a **Series**
 - Methods for **DataFrame**
 - Methods for **Series**
-

So pandas. Its conventionally imported as “pd” for short.

Construct from a list or numpy array:

```
>>> s1 = pd.Series([1,3,5,7,11],name="primes")
```

```
>>> s1
```

```
0    1
1    3
2    5
3    7
4   11
```

```
Name: primes, dtype: int64
```

```
>>> s2 = pd.Series(np.arange(5),name="counter")
```

```
>>> s2
```

```
0    0
1    1
2    2
3    3
4    4
```

```
Name: counter, dtype: int64
```

A Series is like a vector in R, but it has a name attached to it (that's separate from the name of the object). You can construct them from lists. They also have a data type. These here are two numeric integer series.

Construct from Series

```
>>> pd.concat([s1, s2],axis=0)
0      1
1      3
2      5
3      7
4     11
0      0
1      1
2      2
3      3
4      4
dtype: int64
```

```
>>> pd.concat([s1, s2],axis=1)
      primes  counter
0          1         0
1          3         1
2          5         2
3          7         3
4         11         4
```

You can make a Data Frame from Series objects, for example by “concat”, but make sure you get the axis right.

DataFrame

Construct by row from list-of-lists:

```
>>> df = pd.DataFrame([ [1977, "This"], [1983, "That"], [1999,"Other"]])
>>> df
   0    1
0  1977  This
1  1983  That
2  1999  Other

>>> df = pd.DataFrame([ [1977, "This"], [1983, "That"], [1999,"Other"]],
...     columns=["Year","Thing"])
>>> df
   Year  Thing
0  1977   This
1  1983   That
2  1999  Other
```

You can also construct a data frame by rows from the elements in a list. You can name the columns with a "columns" argument.

DataFrame by columns from dict

```
>>> data_dict = dict(Year=[1977, 1983, 1999], Thing=["This","That","Other"])
```

```
>>> data_dict  
{'Year': [1977, 1983, 1999], 'Thing': ['This', 'That', 'Other']}
```

```
>>> pd.DataFrame(data_dict)  
   Year  Thing  
0  1977   This  
1  1983   That  
2  1999  Other
```

If you want to create by columns, the easiest way is probably with a dictionary, where it gets the names from the keys of the dictionary.

DataFrame from a CSV file

```
>>> llc = pd.read_csv("casecontrol.csv")
>>> llc
```

	x	y	cc
0	359014	416976	0
1	352909	426935	0
2	353848	422172	0
3	359202	417326	0
4	357795	415825	0
..
969	356340	413295	1
970	355903	413619	1
971	355563	414116	1
972	355398	414390	1
973	355350	414031	1

```
[974 rows x 3 columns]
```

```
>>> llc.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 974 entries, 0 to 973
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0    x      974 non-null    int64
1    y      974 non-null    int64
2    cc      974 non-null    int64
dtypes: int64(3)
memory usage: 23.0 KB
```

Pandas has a `read_csv` method for reading directly from a CSV file. You can also change the dialect with some other arguments.

Once you've got a pandas data frame, you can get a quick summary of the columns with the "info" method.

DataFrame methods

https://pandas.pydata.org/docs/getting_started/comparison/comparison_with_r.html

Querying, filtering, sampling

R	pandas
<code>dim(df)</code>	<code>df.shape</code>
<code>head(df)</code>	<code>df.head()</code>
<code>slice(df, 1:10)</code>	<code>df.iloc[:9]</code>
<code>filter(df, col1 == 1, col2 == 1)</code>	<code>df.query('col1 == 1 & col2 == 1')</code>
<code>df[df\$col1 == 1 & df\$col2 == 1,]</code>	<code>df[(df.col1 == 1) & (df.col2 == 1)]</code>

There's a guide to Pandas for R users online that goes through all the differences and shows corresponding code for pandas and R data frames.

Basics

Subset with a name = a **Series**

```
>>> llc["x"]
0      359014
1      352909
2      353848
3      359202
4      357795
...
969    356340
970    355903
971    355563
972    355398
973    355350
Name: x, Length: 974, dtype: int64
```

Subset with an array = **DataFrame**

```
>>> llc[["x"]]
      x
0      359014
1      352909
2      353848
3      359202
4      357795
...
969    356340
970    355903
971    355563
972    355398
973    355350
[974 rows x 1 columns]
```

For example, if you subset with a name you get the column with that name as a Series. If you subset with a list you get a data frame. Note that although this looks like R's double-square-bracket method for getting list elements, what you are seeing is a list, constructed by the inner square brackets, inside the square brackets subsetting the data frame.

More basics

Subset with list of columns:

```
>>> XY = ["x", "y"]
```

```
>>> llc[XY]
```

	x	y
0	359014	416976
1	352909	426935
2	353848	422172
3	359202	417326
4	357795	415825
...
969	356340	413295
970	355903	413619
971	355563	414116
972	355398	414390
973	355350	414031

```
[974 rows x 2 columns]
```

Here's that in two stages. XY is a list with two names, and when I subset I get the data frame with those columns.

Basics

Get Series with .column

```
>>> llc.x
0      359014
1      352909
2      353848
3      359202
4      357795
...
969    356340
970    355903
971    355563
972    355398
973    355350
Name: x, Length: 974, dtype: int64
```

Create Series with index

```
>>> llc["x0"] = llc.x - 355000
```

You can also extract a series using dotted notation. You can't create a new series like this though, you have to do it via a named subset.

Index with a slice

Get rows:

```
>>> llc[:3]
```

	x	y	cc
0	359014	416976	0
1	352909	426935	0
2	353848	422172	0

If you subset with a slice, you get rows.

Categorical Series

Change existing column to categorical:

```
>>> llc.cc = pd.Categorical(llc.cc)
>>> llc
```

	x	y	cc	x0
0	359014	416976	0	4014
1	352909	426935	0	-2091
2	353848	422172	0	-1152
3	359202	417326	0	4202
4	357795	415825	0	2795
..
969	356340	413295	1	1340
970	355903	413619	1	903
971	355563	414116	1	563
972	355398	414390	1	398
973	355350	414031	1	350

Pandas' equivalent to R's factors is the Categorical data type. You can create a categorical series from any other series, and it sets the unique values.

Change Category Names

Modify the `.cat`
attribute of the
Series

```
>>> llc.cc.cat.categories = ["Control", "Case"]
>>> llc
```

	x	y	cc	x0
0	359014	416976	Control	4014
1	352909	426935	Control	-2091
2	353848	422172	Control	-1152
3	359202	417326	Control	4202
4	357795	415825	Control	2795
...
969	356340	413295	Case	1340
970	355903	413619	Case	903
971	355563	414116	Case	563
972	355398	414390	Case	398
973	355350	414031	Case	350


```
[974 rows x 4 columns]
```

The labels of a categorical series can be changed, this is like changing the levels of a factor in R.

Table/Plot of categories

```
>>> llc.cc.value_counts()  
Control    917  
Case        57  
Name: cc, dtype: int64
```

```
>>> llc.cc.value_counts().plot(kind="bar")  
<matplotlib.axes._subplots.AxesSubplot object at 0x7ff4d391d1c0>
```



You can get basic statistics of a category and a simple barplot.

Summaries via `.describe()`

```
>>> llc.describe()
              x              y              x0
count      974.000000      974.000000      974.000000
mean    355525.837782  421518.232033    525.837782
std       3365.459162    4527.043338    3365.459162
min     346475.000000  412437.000000 -8525.000000
25%     353031.000000  417358.250000 -1969.000000
50%     355869.500000  421899.500000    869.500000
75%     358202.000000  425984.500000    3202.000000
max     364435.000000  428987.000000    9435.000000
>>> llc.describe(include="category")
              cc
count          974
unique          2
top      Control
freq          917
```

The equivalent of “summary” is the “describe” method which gives summary statistics of each numeric column. If you want info about the categorical columns you can ask for those.

All the things...

```
>>> llc.describe(include="all")
```

	x	y	cc	x0
count	974.000000	974.000000	974	974.000000
unique	NaN	NaN	2	NaN
top	NaN	NaN	Control	NaN
freq	NaN	NaN	917	NaN
mean	355525.837782	421518.232033	NaN	525.837782
std	3365.459162	4527.043338	NaN	3365.459162
min	346475.000000	412437.000000	NaN	-8525.000000
25%	353031.000000	417358.250000	NaN	-1969.000000
50%	355869.500000	421899.500000	NaN	869.500000
75%	358202.000000	425984.500000	NaN	3202.000000
max	364435.000000	428987.000000	NaN	9435.000000

If you ask for everything you get a table with a bunch of missing elements where it doesn't make sense to measure that for that type of column!

Grouped Aggregates

```
>>> llc.groupby("cc").agg({'x': 'mean', 'y': 'mean'})
              x              y
cc
Control  355566.651036  421514.501636
Case     354869.245614  421578.245614

>>> from statistics import median
>>> llc.groupby("cc").agg({'x': median, 'y': median})
              x              y
cc
Control  355980  421890
Case     355207  422333
```

There are methods for grouping and computing aggregates over groups too. Here I group over the CC column and compute the mean x and y values. I can use any function here, such as “median” imported from Python’s statistics module.

Pickling Python Data

```
>>> import pickle
```

```
>>> data_list = [{"Name","Age"}, ["Al",22], ["Bob",44], ["Chip",33]]
```

```
>>>
```

```
>>> pickle.dump(data_list, open("data_list.pkl","wb"))
```

Write in binary mode

Write this object to
this file connection

```
>>> del data_list
```

```
>>>
```

```
>>> pickle.load(open("data_list.pkl","rb"))
```

Read back from this
connection

```
[['Name', 'Age'], ['Al', 22], ['Bob', 44], ['Chip', 33]]
```

Time for some more general data things in Python. Pickling. If you want to save an object in your Python session and its too complex for a CSV file, you can “pickle” it. Use the “pickle” module and dump it to a file. Make sure the file is opened in “wb” mode for writing binary data. Then you can get your object back by loading it.

Structured Data

- JSON
- XML
- YAML

The pickle format is specific to Python, so if you want to interchange data with other systems, such as R for example, you need a data interchange format. There are plenty of them, I'll look at three here.

JavaScript Object Notation

```
>>> import json
>>> s = { 'name': 'GOOG', 'shares': [100,110,120], 'price':490.1 }
>>>
>>> s
{'name': 'GOOG', 'shares': [100, 110, 120], 'price': 490.1}
>>>
>>> json.dumps(s)
'{"name": "GOOG", "shares": [100, 110, 120], "price": 490.1}'
```

JSON is a text format for structured data. There's a built-in module for it in Python. You can get a JSON representation of a Python object as a string using the "dumps" method, which you could then write to a file. Notice how its very similar to - but not quite exactly the same as - the way Python prints its objects out. JSON is very popular for data interchange and there's packages for most common modern programming languages.

XML : eXtensible Markup Language

Highly Structured, Formalised

```
<?xml version="1.0"?>
  <data date="2021-12-02">
    <variables>
      <a>100</a>
      <b>Hello</b>
      <z max="1000" min="0">666</z>
    </variables>
  </data>
```

The big beast of data interchange formats is XML. This is data interchange designed by committee, and is often referred to as an “enterprise” format. Here’s a very simple example, but the XML standard specifies all sorts of complexities, including methods for transforming from one XML format to another expressed in XML. It can do anything. The problem is that its not really for human consumption.

Human and Computer-readable

```
quiz:
  description: |
    This is another quiz, which
    is the advanced version of the previous one
  questions:
    q1:
      desc: "Which value is no value?"
      ans: Null
    q2:
      desc: "What is the value of Pi?"
      ans: 3.1415
```

In response to XMLs complexity, in cases where less complexity is needed, YAML (yet another markup language) was developed. This removes all the angle brackets and tags for a nested structure defined by indenting, and a tag: value system. This is readable enough to be human readable and writable, but also structured enough for computers to read it. And modules exist for reading YAML for most of the common modern languages - for Python the thing you need is the pyyaml package.

Databases


```
>>> con = sqlite3.connect("sample.gpkg")
>>>
>>> query = con.execute("select time_start, rainfall from annual limit 10")
>>>
>>> rain = query.fetchall()
>>> rain = pd.DataFrame(rain)
>>> rain
```

	0	1
0	1784-01-01	35.12
1	1785-01-01	36.69
2	1786-01-01	32.25
3	1787-01-01	51.01
4	1788-01-01	29.36
5	1789-01-01	55.88

Lots of data is kept in organised databases. The way to access these is similar to accessing files - you open a connection, do stuff with the connection, and then close the connection to tidy up. Here I'm connecting to a database stored in a file called "sample.gpkg", which is an SQLite3 file. Then I can run access data by querying tables in that database and converting them to pandas data frames.

Databases

```
>>> query = con.execute("select time_start, rainfall from annual limit 10")
```



Connection could be to a database on disk, a local server, or a network cloud service.

One of the advantages of the database system is that once you have the connection, your code is the same whether you have a connection to a local file, a networked local database, or some cloud data service. As long as there's a "driver" for your data source type, you can do this.

Straight to Pandas

```
>>> con = sqlite3.connect("sample.gpkg")
>>> rain = pd.read_sql("SELECT time_start, rainfall FROM annual", con)
>>> rain
```

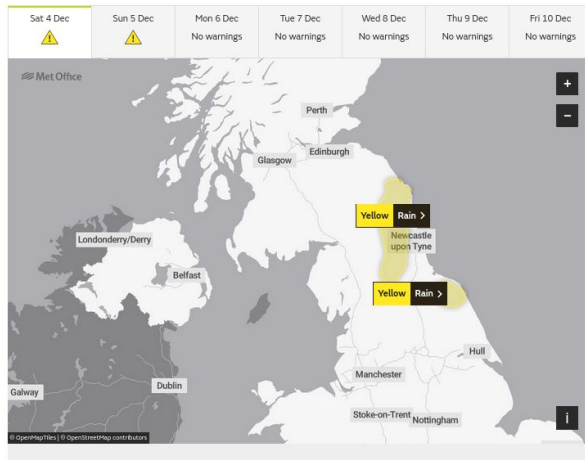
	time_start	rainfall
0	1784-01-01	35.12
1	1785-01-01	36.69
2	1786-01-01	32.25
3	1787-01-01	51.01
4	1788-01-01	29.36
..
363	1915-01-01	41.16
364	1916-01-01	49.28
365	1917-01-01	41.97
366	1918-01-01	51.95
367	1919-01-01	39.71

```
[368 rows x 2 columns]
```

And you can go straight to pandas data frames with some pandas methods once you have a connection.

Web Scrapping

UK weather warnings



```
<!DOCTYPE html>
<html lang="en">
<head>
<link rel="preconnect" href="https://cc.cdn.civiconputing.com" crossorigin="anonymous">
<link rel="preload" href="/webfiles/1637839662685/js/common/set-body-class.js" as="script">
<link rel="preload" href="/webfiles/1637839662685/js/common/run-ab-test.js" as="script">
<meta charset="utf-8"/>
<meta http-equiv="X-UA-Compatible" content="IE=Edge">
<meta name="viewport" content="initial-scale=1.0, width=device-width, viewport-fit=cover">
<meta name="eGMS.accessibility" content="A" />
<meta name="eGMS.subject.category" content="Climate and weather" />
<meta name="twitter:site" content="@metoffice" />
<meta name="twitter:card" content="summary" />
<meta property="fb:app_id" content="208580322593329" />
<meta property="og:site_name" content="Met Office" />
<meta property="og:type" content="website" />
<meta property="og:url" content="https://www.metoffice.gov.uk/weather/warnings-and-advice/uk-warn" />
<meta name="theme-color" content="#2a2a2a" />
<meta property="og:image" content="https://www.metoffice.gov.uk/binaries/content/gallery/metoffice/ih_lnxvxlqrdqb2kyrixw3lmmul7n-abr2juzvzvz07_8mk04g7wuma15ec6yt11de42bgxoxdizc75bhesr6ywacs9tbyrery9" />
<meta property="og:image:width" content="1596" />
<meta property="og:image:height" content="880" />
<meta name="google-site-verification" content="LR7Q3483mkrmLc3J2pzTbL48MnTbBDVzjA3GqT3Q2W"/>
<meta name="google-translate-customization" content="f8c19e614cd5ee7e-37428b41ac444aca-gd1bc91f73" />
<link rel="canonical" href="https://www.metoffice.gov.uk/weather/warnings-and-advice/uk-warnings" />
<link rel="stylesheet" href="/webfiles/1637839662685/css/compiled/base.css" type="text/css"/>
<script src="/public/adhoc/no-adhoc.js" defer=></script>
<script src="/public/gdpr-consent/cookieControl-config.js" defer=></script>
<link rel="stylesheet" href="/webfiles/1637839662685/css/compiled/components/breadcrumb.css" type="text/css"/>
<link rel="stylesheet" href="/webfiles/1637839662685/css/compiled/components/feedback-banner.css" type="text/css"/>
<title>UK weather warnings - Met Office</title>
<meta property="og:title" content="UK weather warnings"/>
<meta name="title" data-type="string" class="sui-type" content="UK weather warnings"/>
<meta name="description" content="Met Office UK weather warnings for rain, snow, wind, fog and ice to date with local weather warnings."/>
<meta property="og:description" content="Met Office UK weather warnings for rain, snow, wind, fog and ice to date with local weather warnings."/>
```

Another common source of data, some of it informal and of possible dubious authenticity or legality, is web pages. Pretty much anything you see on a web page is on your computer (otherwise how would you see it?) and so could potentially be converted into useful data. This process is known as “web scraping”. This is the Met Office UK Weather Warnings page and shows here a couple of rain warnings. I can look at the web page source and I know that somewhere in here must be the data that defines those warnings and the warning text and colour and so on.

Scrape scrape...

Line 1731

```
<script>
//
window.metoffice = window.metoffice || {};
window.metoffice.warnings = window.metoffice.warnings || {};
window.metoffice.warnings.geojson = {"type":"FeatureCollection",
  "properties":{"weatherType":["rain"],"validFromDate":"2021-12-04T",
  "headline":"Persistent, possibly heavy rain, with some snow on t
  ime travel disruption.", "validToDate":"2021-12-05T05:00Z", "whatI
  roads, and snow on higher level routes, probably making journey
  ices probably affected with journey times taking longer", "Floc
  is is possible"}, "warningId":"0d9246a7-bfd6-4274-blea-b04254ea22
  i":"2021-12-05"}}, "geometry":{"type":"MultiPolygon", "coordinates":
  i37], [-1.0052, 54.5434], [-1.1041, 54.5019], [-1.2634, 54.4381], [-1.1
  i1.0657, 54.3133], [-1.0052, 54.2748], [-0.7581, 54.2492], [-0.5823, 54
  i, 54.4732], [-0.769, 54.4988]]], [[[-1.615, 55.5659], [-1.6919, 55.562
  i745], [-2.0929, 55.6373], [-2.0819, 55.5286], [-2.1973, 55.4197], [-2.
  i-2.2522, 55.2635], [-2.0984, 54.8482], [-2.1735, 54.7257], [-2.195, 54
  i54.5593], [-1.7798, 54.6579], [-1.7139, 54.7056], [-1.7084, 54.7817],
  i6], [-1.6809, 55.2322], [-1.637, 55.2791], [-1.6479, 55.4726], [-1.626
  i];
  //]]&gt;
&lt;/script&gt;</pre></div><div data-bbox="591 165 922 434" data-label="Figure"><img alt="Map of Great Britain and Ireland showing weather warnings for rain and snow. The map highlights areas in Scotland, Northern Ireland, and parts of England and Wales. Specific dates like 2021-11-26 and 2021-11-27 are marked on the map, indicating the duration of the warnings. Major cities like London, Manchester, and Birmingham are labeled."/>A map of Great Britain and Ireland with weather warnings overlaid. The warnings are represented by colored polygons: orange for rain and red for snow. The map shows warnings for rain and snow across various regions, including Scotland, Northern Ireland, and parts of England and Wales. Specific dates like 2021-11-26 and 2021-11-27 are marked on the map, indicating the duration of the warnings. Major cities like London, Manchester, and Birmingham are labeled.</div><div data-bbox="105 482 875 558" data-label="Text"><p>And I find it at line 1731 inside a <code>&lt;script&gt;</code> tag as part of a “geoJSON” object. Its then not too hard for me to extract that geoJSON and load it into my favourite geographic information system and map it (this is a different set of warnings from the previous page).</p></div>
```

Application Programming Interface



<http://api.geonames.org/timezoneJSON?lat=47.01&lng=10.2&username=demo>

hostname endpoint query parameters username/pass

To save people having to scrape data from web pages and encourage a more formal approach to getting useful data, some companies expose an API for their data service. This is a URL that will get you some timezone info for this lat-long coordinate from a service called “Geonames”. If you go to this URL in a browser (and you may need to register a username first and replace it here) you get back some JSON data:

Application Programming Interface

<http://api.geonames.org/timezoneJSON?lat=47.01&lng=10.2&username=demo>

```
{
  "countryCode" : "AT",
  "countryName" : "Austria",
  "dstOffset" : 2,
  "gmtOffset" : 1,
  "lat" : 47.01,
  "lng" : 10.2,
  "rawOffset" : 1,
  "sunrise" : "2021-12-04 07:47",
  "sunset" : "2021-12-04 16:30",
  "time" : "2021-12-04 16:15",
  "timezoneId" : "Europe/Vienna"
}
```

```
>>> import geocoder
>>> g = geocoder.geonames('New York', 1)
>>> g.address
'New York City'
>>> g.geonames_id
5128581
>>> g.description
'city, village, ...'
>>> g.population
8175133
```

And that data looks like this - the point is in Austria and the timezone and current time is given, together with some other data. Now you could write some Python code to do all this - there's a builtin module for doing web requests in Python - but for many popular services a quick search on PyPI will find that someone has already written a complete module for you. Here's the "geocoder" module with has a method for calling the "geonames" web services and getting info back easily.

Summary

- Data is? Data are?
 - Reading/Writing files in Python
 - Pandas for Data Frames
 - Getting Data from Sources
-

I've not answered the important question - is "data" singular or plural? But I've outlined basic data read/write operations in Python, and given an outline to Pandas for data frame work, and shown a bit about sources of data that a modern data scientist is going to encounter. Although people will still email you Excel spreadsheets and data in tables in PDFs...